

EXTENDED PROJECT QUALIFICATION

To What Extent Will Sectors Change with the Mass Introduction of AI?

Candidate Name:

Sebastian Cawthorne

Centre Number:

40404

Candidate Number:

6023

Word count:

5200

Contents

Title page	1
Contents	2
Introduction	3
What is AI?	3
Prediction One: How Employment Will Change	4
Prediction Two: A Sector Deep-Dive into Military and Defence	6
Prediction Three: Skill Sets Will Need to Adapt	8
Prediction Four: How the Education System Must Adapt	10
Conclusion	12
Bibliography	14
Appendix	16

Artificial intelligence (AI) is one of the most significant technological developments of the twenty-first century, and it is changing the structure of economies, workplaces, and education systems faster than most people realise. In this essay I am going to look at the question: to what extent should sectors change with the mass introduction of AI? To do this I have drawn on academic sources, industry research, and a survey I carried out with 57 professionals across multiple sectors, as well as my own experience working at Databricks and Microsoft. My overall argument is that whilst AI will affect every sector to some degree, how much each one needs to change, how fast, and in what way will depend heavily on the sector, the role, and how ready people are to adapt.

Despite how widely it is used, AI is still a very misunderstood topic. It is not simply a search engine returning the closest answer to your question; it is far more sophisticated than that. Rossi (2018) argues that before society can fully trust AI, people first need to understand what it actually is and how it works. With that in mind, I want to start by breaking AI down into three core types of AI, with AI being the general term for all codes that produce human - like thinking. The subsections (Toloka Team, 2023):

Machine learning – the field of AI where computers learn patterns from data and improve their performance on a task without being explicitly programmed with all the rules.

Large language models (LLMs) – this is what most people will use in day-to-day life. It is trained on huge amounts of data like the internet and is built around the idea of understanding and generating human language. Some examples of this include ChatGPT which is a large language model built to help humans.

Generative AI – this is typically built for new content instead of gathering data from the internet. It is able to generate audio, photos, videos and more. This is usually integrated in AI systems like Grok which has photo creation and GPT which can speak back to you acting like a real person. Most large AI systems like ChatGPT have both a large language model and generative AI integrated into them.

From now on I will be referring to it as AI (most of what I will be talking about will be generative AI and large language models as that is the type almost everyone will use in their day-to-day life). All AI types, except machine learning, are relying on data accessed via the internet and internal corporate systems. The problem we are now running into is that now AI is creating and republishing data from the internet, it is increasingly using data gathered and generated by

itself, which can create a degradation of the AI model over generations of training. This leads to incorrect data and is called "model collapse". In addition, AI can "hallucinate" every now and again, (meaning it makes up false information or uses false information). Model collapse and hallucinations can cause a dramatic increase in the amount of false information produced by AI every time it is used. This is why it is so important that sources are checked carefully, as AI may be producing unreliable answers. In my view, this is one of the most important things to keep in mind when thinking about how fast we should be adopting AI (Rossi, 2018).

In this section of my EPQ I will be discussing which sectors are already being affected by AI and which will be in the future. As AI is evolving so fast, we can already see some huge changes that have happened and this aids us, alongside research in this field, to make a well-educated prediction on the future and how AI will shape the market. My predictions below are based on a range of research, a survey I carried out with 57 professionals across different industries, and my own experience working at Databricks and Microsoft, which gave me a real insight into how AI is being used in practice.

Prediction One: How Employment Will Change

From my research I can see that AI will make a huge number of roles far more efficient. For example, researchers will be able to gather and look through much larger volumes of data much more quickly, which will lead to more accurate results (Jung and Desikan, 2024). On top of that, I think AI has the potential to massively increase how much work one person can get done, if they are trained to use it properly. This came through clearly in my own survey, where 78% of respondents agreed that companies should train their staff to use AI. I also came across evidence that AI can do certain tasks more accurately than a person working on their own, especially in healthcare diagnostics, where AI systems have beaten individual doctors in controlled tests. (Sharps, 2024)

I see every single sector being affected by AI, but some sectors and roles being affected faster. My primary research agrees, as when I asked which sectors they thought would be most transformed by AI, healthcare came out on top with 67% selecting it, followed by transportation and logistics at 47% and financial services at 42%. This shows that people don't think they will be affected at the same rate. Roles will start with repetitive admin jobs; examples include secretaries, call centre agents, paralegal-type work, compliance checks for various industries, as just a few examples. Personally, I can see almost all of these jobs being replaced by AI within the next 10 years at the latest. I can also see the requirement for some roles to be in person to make sure everything is running smoothly, be that on the subject matter at hand or ensuring the

AI is delivering trusted results. Overall, there will be a dramatic decrease in employees in these roles.

These are the roles I see being hit first and hardest. That said, I do not think the story of AI is simply about replacing people. In healthcare for example, I can see AI-assisted transcription and diagnostic tools taking the pressure off already stretched services and helping to plug skills shortages, without necessarily cutting staff numbers. (Sharps, 2024) To me, the important question is whether AI is replacing workers or just helping them do more. In journalism I see something similar: AI will probably increase how much one journalist can produce rather than wiping out the whole profession, although you would likely need fewer people to produce the same amount of content (Jung and Desikan, 2024) which will either lead to reduced headcount at a firm or an increase in the amount of content released from the same size firm.

While I've outlined plenty of opportunity, we must also weigh a key barrier: people's reluctance to adopt AI, which could slow the pace of job change. Some of the problems we might face in automating these roles are that people value human interaction, and because AI is still so new, many remain distrustful of it. This is a very real concern in my experience - AI doesn't always get things right first time. I therefore suspect that many people in highly repetitive jobs may struggle to adopt it, slowing its uptake in these roles and delaying the point at which employment in them begins to decline.

To explore what the future might look like, I will now turn to a research paper from Goldman Sachs (Briggs and Dong, 2025). Their research suggests that AI will affect jobs, pointing in particular to the gains in efficiency it offers for our work. Goldman Sachs predicts that for every 1 percentage point gained in productivity via technology there will be a 0.3% increase in unemployment. On its own that might not sound like much, but because AI keeps improving itself, the impact is going to build up over time. If productivity goes up at just 1% a year, which I think is a pretty conservative estimate, that works out to roughly 3% more people unemployed every decade. If this pattern holds, the result could be structural unemployment at scale within a single generation. There will also be an increase of about 3% unemployment every decade. This was also reflected in my survey where when I asked whether AI would significantly reduce jobs within the next ten years. My respondents averaged a score of 3.3 out of 5, which I think shows people are genuinely concerned but not panicking. 42% leaned towards agreeing and not a single person thought no jobs would be affected at all, which aligns with what Goldman Sachs are predicting. Another prediction from Goldman Sachs is that there is a small chance that there will be an opportunity for the job market to grow significantly. Some previous

evidence that they used was the creation of the internet. 60% of all jobs in the US did not exist in the 1940's and 85% of these are now technology driven. I also think that the industrial revolution is another example of new technology removing a lot of jobs but still making more than it displaced at the time due new roles being created. Based on these previous examples, I strongly agree with Goldman Sachs view that there will most likely be a period of higher unemployment during the adoption of AI. However, this adoption period could be much shorter or much longer, as I think more people are becoming sceptical of AI given that humans have less control over it as it develops. Another factor that might slow the adoption of AI is the influence of sci-fi films depicting AI going rogue: The Matrix, Blade Runner and The Terminator just to name a few. I think these films could cause significant hesitation around AI especially when technology companies want to start using it in robots. However, many companies will help to speed up the adoption of AI due to their desire to reduce costs as much as possible. What this means in practice is that people who can use AI effectively and responsibly will have a real edge when applying for jobs, due to them being able to produce more work in the same amount of time. I will look at this in more detail later on. However, it is worth noting that this Goldman Sachs data does not cover much of the world, as it reflects a largely Western perspective. Many people in developing countries will not have widespread access to AI for years to come. The unemployment data is also focused on the US, rather than Europe or other high-income countries.

Prediction Two: A Sector Deep-Dive into Military and Defence

There are a lot of sectors that will be affected by AI and all of them will have some level of controversy around them, mostly around unemployment. But I think the ones that will cause the biggest controversy are the ones where human lives are directly on the line. That includes the military, surgeons, healthcare and transport. What all of these have in common is that a mistake does not just mean a bad business outcome, it means someone could die. I think this is why they will probably be some of the last sectors to have AI fully introduced at scale, and also why they will attract the most debate when it does happen.

Starting with the military, I think there are three main ways AI is going to change things. The first and probably the least controversial is logistics. Right now, a huge amount of military resource goes into just making sure the right equipment is in the right place at the right time. I think AI could take a lot of that work off people's plates and do it more accurately than a human could, a bit like a just-in-time system but without the human error. That frees up more time and

attention for actual strategy and decision-making, and it would probably save money too by reducing waste and unnecessary storage.

The second area is research and development. AI is already being used to run simulations and test ideas before anyone builds a physical prototype, and I think that is only going to increase. It massively cuts down the time and cost of developing new technology, and whichever countries adopt this first are going to have a real advantage over the ones that don't (Bondar, 2024).

The third area is the one I find most worrying, which is autonomous weapons and drones. Unlike the other two, this is not just about efficiency. This is about whether we are comfortable letting an AI system make the decision to take someone's life with no human involved. People sometimes bring up Asimov's three laws of robotics here, the idea that robots shouldn't harm humans, should follow orders, and should protect themselves. The problem is that those aren't actual laws. They came from a science fiction novel written in the 1950s and have no legal standing anywhere. So right now, there is genuinely no binding international rule stopping a government or company from building a weapons system that kills people on its own judgement.

The UN Secretary-General flagged this in 2024 and called for a proper legal framework to be in place by 2026 specifically to ban autonomous weapons that operate without human control (United Nations, 2024). That framework still doesn't exist, which I think is a real problem. We have already seen major AI companies in the US being asked whether their systems could be used to make the decision to fire a missile. They all said no, but the fact that there is nothing legally stopping them from saying yes is concerning.

My biggest worry is that AI just isn't reliable or empathetic enough yet to be making these kinds of decisions. It works off statistics and probabilities, not human judgement, and I think that massively increases the risk of civilians being misidentified and killed (Toh and Ayoub, 2026). Someone in my survey put it really well, saying that AI in military contexts risks "reducing human life to a numerical value and acting on probabilities rather than protection of life." I have already seen this starting to play out with reports of the US using AI systems like Maven Smart to generate targeting recommendations in Iran, pulling data from satellites, drones, social media and data brokers to flag people and objects of interest. A human still has to approve the final call, which is something, but I still think it is worrying because if the AI is fed bad data, or if the person approving just rubber-stamps what the AI says, things could go very wrong very quickly. Drones are another huge area where this is developing fast, as anyone watching what has been

happening in Ukraine can see. It does make you wonder whether we need proper international laws around this before it goes any further.

Prediction Three: Skill Sets Will Need to Adapt

Something I have been thinking about a lot is whether the skills needed for most jobs have actually changed yet, or whether AI is just helping people do the same things a bit faster. From what I have found, the answer is different depending on the type of job. The picture is a split one. For primary and secondary sector roles – typically lower-paid, more manual positions – the core requirements of commitment, reliability, and physical capability are unlikely to shift dramatically in the near term. However, for tertiary and quaternary sector roles, the shift is already underway. In my own survey, where I asked respondents to rank six workplace skills by future importance, communication and interpersonal skills were consistently ranked most critical, followed by critical thinking, with the effective use of AI tools ranked third (Primary Survey, 2026). This reflects a broader trend identified by Sharps (2024): as AI increasingly handles information retrieval and first-draft generation, the premium placed on raw knowledge retention is declining, whilst the ability to interrogate, interpret, and communicate AI outputs is rising. As one Finance Director surveyed for this essay noted: "The most valuable people will be those who can ask good questions, spot weak answers, communicate clearly, adapt quickly and use human judgement." I have also noticed that more students seem to be getting higher grades at university, which I think is partly down to the growing use of AI in academic work, meaning people don't stand out as easily when applying for jobs. I now think AI is getting good enough to do most of the work for the average person. You then just need to articulate what the AI has produced so that the other person you're communicating with can understand it properly.

Something else I think is worth considering is how this shift in skills is affecting different age groups differently. Older workers who have spent years building up knowledge in their field now find that AI can replicate a lot of that knowledge almost instantly, which I think must be quite unsettling. However I also think younger people entering the workforce face a different but equally difficult challenge. They are coming into a job market where the bar for what counts as a basic level of competence is rising fast because AI is raising the floor for everyone. Sharps (2024) makes the point that as AI handles more of the routine cognitive work, the gap between

someone who uses it well and someone who does not is going to widen significantly over time. I think this is actually one of the most important things young people need to understand right now. It is not enough to just use AI, you need to use it better than the person sitting next to you. On top of that I think there is a real risk that people who rely on AI too heavily early in their careers never actually build the underlying judgement that makes you good at a job in the first place. The best people I met at Microsoft and Databricks had years of experience that gave them the instinct to know when the AI was right and when it needed challenging, and that is not something you can shortcut.

I also think companies are starting to believe that it is relatively easy to teach someone a technical skill, but much harder to teach the soft skills that are aligned to personality traits. I think employers will increasingly look for people who can use AI effectively and responsibly to get the most out of their work, and who are keen to learn and evolve alongside the technology. This is likely what they would prioritise if a candidate lacked specific knowledge or experience.

My work experience at Microsoft and Databricks in April and May 2026 gave me a real insight into how AI is being used in practice inside two major tech companies. Over five days working within the enterprise sales teams at both companies, I sat in on internal meetings, shadowed staff at all levels, and took part in live sales exercises with feedback from working account executives. The thing that stood out to me most was that the people who were best at their jobs were not the ones who knew the most about AI technically (they were never going to be as technical as a solution architect however some were more up to date than others). They were the ones who were best with people. Being able to build a relationship, ask the right questions, and explain something complicated clearly to a customer counted for far more than knowing every detail of how the product worked under the bonnet. This backed up what my survey found, and honestly it confirmed something I had suspected all along: even at companies that are building and selling AI, the people skills are what make the difference.

The other thing I noticed was how much AI had already changed what the sales job actually involves on a day-to-day basis. People I spoke to told me that a lot of the prospecting work, things like researching leads and putting together account information, is now done by AI tools rather than by the salesperson themselves. I was also shown a demo of how AI was being used to sort through customer complaints, group similar ones together, and pull out a summary so a human could deal with them much more quickly. The human was still the one making the decisions in both situations, the AI was just doing the prep work beforehand. Seeing this in real

life gave me a much clearer picture of what I have been arguing about throughout this essay: that AI is changing what jobs look like rather than just getting rid of them.

Prediction Four: How the Education System Must Adapt

Because of all the changes to the job market and the skills employers are now looking for, I think schools and universities urgently need to catch up. From everything I have read, it feels like the education system is still preparing students for a world that is already being left behind. The biggest change I want to see is schools stopping the blanket ban of AI and starting to actually teach students how to use it properly. Chen (2025) makes a good point that even subjects like humanities need to engage with AI rather than just blocking it out. I think the comparison to typing lessons is a really useful one here: when computers came in, schools adapted and started teaching students to type because it was such a useful skill. I think AI prompt writing should be treated the same way. Cutler (2025) agrees that AI can genuinely improve how subjects are taught if it is properly integrated. In my own survey of teachers, most said they supported teaching students to use AI responsibly, and the ones who were cautious were mainly worried about students losing the ability to think for themselves rather than being against AI in principle.

One argument schools make that I do think has some merit is that understanding how something works matters more than just knowing the answer it produces. If a student uses AI to get the answer to a maths problem without understanding the method, they will struggle the moment AI is not available or gives a wrong answer and they cannot spot it. The same applies across most subjects. A doctor who relies on AI diagnostics without understanding the medicine behind them is a dangerous doctor, and a lawyer who uses AI to draft a contract without understanding the law is a dangerous lawyer. There is a genuine case for saying that AI should be a tool you apply on top of knowledge, not a way of avoiding building it in the first place.

That argument makes sense up to a point, but I think it misses what is actually changing. The body of knowledge required to do most jobs well is growing so fast that no person can keep up with all of it on their own anymore. AI does not remove the need to understand things; it changes what level of understanding is actually necessary. A doctor does not need to memorise every drug interaction if an AI system can cross-reference it instantly and more accurately than human memory ever could. What that doctor needs are the judgement to know when to trust the output and when to question it. That is a different skill to memorisation, and I would argue it is a harder one. The Education Endowment Foundation ranks metacognition,

which is the ability to think about and evaluate your own thinking, as one of the highest impact and lowest cost interventions in education, adding an average of seven months of additional progress for students. (Churches and McAleavy, 2019) Teaching students to interrogate AI outputs, spot errors, understand the limits of a model and ask better questions is exactly this kind of thinking. It is a deeper form of understanding than being able to recall a fact, and it is the skill that will actually matter in the workplace. My own survey backed this up, with critical thinking ranking second out of six skills when I asked professionals what they value most in young people, above technical knowledge and above the ability to use AI tools themselves.

The main worry, which comes up a lot in the research I looked at, is that students might stop actually thinking for themselves if they just let AI do everything for them. Johnson-Roehr (2023) calls this cognitive offloading, and there is evidence that it can genuinely weaken analytical skills over time. But I think the answer to that is not to ban AI, it is to teach students from a young age to use it as a support tool rather than a replacement for their own thinking. If that is done properly, I do not think the younger generation will just become dependent on it.

The situation we are in at the moment is very similar to the introduction of the calculator. When the calculator was first introduced a lot of teachers argued that it would dramatically reduce the mental maths capability of maths students and would hinder their critical thinking skills, and some would also say it was cheating. To some extent it did reduce the mental maths ability for students. Although it meant students stopped doing basic arithmetic in their heads, it also meant they could tackle far harder problems than before. I do not think students got worse at maths overall, they just moved on to thinking about harder things. I think the same will happen with AI. The writing and research skills that teachers worry students are losing might actually just be getting redirected toward higher level thinking like evaluating sources, pulling ideas together, and communicating clearly, which are exactly the skills that came out on top in my own survey when I asked professionals what they actually value in young people entering the workforce.

Now if we look at exams that students take, at least in the UK, there is normally a few exam boards that you choose from and then it is a set exam where everything is pretty much standardised from similar topics and very similar exam papers. Now what do I think should change with AI and modern technology? Well if there were no external facts I had to rely on I would say that the exam boards (apart from a few like medicine should follow these rules) should be done on a computer simply due to the fact that 99% of business with over 10 employees have some form of digitalised data (GOV.UK, 2024) This shows how big and

prominent typing on a computer is and if somebody can get typing at a decent speed they are much more effective and efficient.

A much more controversial change I would like exam boards to make concerns the rule that nobody is allowed to use the internet. The internet is so vast and plays such a role in everyone's life whether they like it or not, that I think we should be able to use it during exams — including AI and any other tools we choose. This would be a positive change because we would be using the same tools we are allowed to use in the real world. It would give a much better representation of how well someone might perform in a company, where employers don't care whether you can remember the year the telephone was invented off the top of your head.

A few problems that would hinder the introduction of AI is that there is some less privileged area where they don't have a computer for everybody. This would cause a lot of problems due to people not having the right equipment to practise for or sit digital examinations. However, in an ideal world I think the whole approach to assessment would need to change, shifting toward real-world problem-solving tasks similar to what you would find in a job assessment centre. On this front, the US system might actually have an advantage over the UK because it does not have the same rigid set of national exams like GCSEs. My overall point is that an education system built almost entirely around memorising facts does not make much sense anymore when AI can look up almost any piece of information in seconds. One idea I find really interesting is testing students on the quality of the prompts they write, because that would actually show whether they understand how AI works and what its limits are, which is a real skill for the modern world (Shoukat, 2024).

In conclusion.

Throughout this essay I have looked at how AI is already changing the world of work, the military, and education, using academic research, a survey of 57 professionals across multiple industries, and my own experience at Microsoft and Databricks. The short answer to the question of how much sectors should change is: a lot, but it will look different depending on which sector you are talking about.

For jobs and employment, the evidence points clearly to repetitive and administrative roles being the most at risk in the short term, with Briggs and Dong at Goldman Sachs suggesting up to 6-7% of the US workforce could be displaced if AI is widely adopted (Briggs and Dong, 2025). My own survey backed this up, with 96% of respondents identifying repetitive admin roles as the most at risk (Primary Survey, 2026). That said, I do not think this means mass unemployment

across all sectors. What I saw at Microsoft and Databricks, and what my primary research suggested, is that AI is more likely to reshape jobs than simply get rid of them. Healthcare and transport came out as the sectors my respondents thought would be most transformed, and my reading of the evidence suggests they are right.

For the military, I think the change needs to be the most carefully managed. AI in logistics and research and development is already happening and is broadly positive. Autonomous weapons are a completely different matter, and I think no government should be deploying them without a human in the loop. The United Nations called for a legally binding framework on this by 2026 and the fact that one does not yet exist is a real problem (United Nations, 2024).

For education, the change is the most urgent and the most overdue. Schools are still built around memorisation at a time when AI makes information instantly accessible to anyone. I think the focus needs to shift towards teaching students to think critically, evaluate AI outputs and ask better questions, which the Education Endowment Foundation identifies as some of the highest impact skills we can develop (Education Development Trust, 2025). AI is not going away and pretending otherwise is not a strategy.

Overall, I think the sectors that change fastest and most thoughtfully will come out ahead. The ones that resist or ignore AI entirely risk being left behind, and the people working in them risk the same.

Bibliography

- Bondar, K. (2024). *Understanding the Military AI Ecosystem of Ukraine*. Centre for Strategic and International Studies/JSTOR. doi:10.2307/resrep64744.
- Brooks, M. (2024). *Is Artificial Intelligence Racist? How Algorithms in the Criminal Legal System Disproportionately Harm People of Colour* [Honours thesis]. Muhlenberg College. doi:10.2307/community.37915972.
- Chen, Z. (2025). The Humanities in an AI World: Personal Reflections from SXSW EDU 2025. *About JSTOR*. Available at: <https://about.jstor.org/blog/the-humanities-in-an-ai-world-personal-reflections-from-sxsw-edu-2025/> [Accessed 12 January 2026].
- Churches, R. and McAleavy, T. (2019). *Metacognition: why teaching students to think about their learning matters more than ever*. [online] Edt.org. Available at: <https://www.edt.org/insights-from-our-work/metacognition-why-teaching-students-to-think-about-their-learning-matters-more-than-ever/> [Accessed 14 May 2026].
- Cutler, C. (2025). AI for schools: Unlocking potential in curriculum processes. *My college*. Available at: https://my.chartered.college/impact_article/ai-for-schools-unlocking-potential-in-curriculum-processes/ [Accessed 16 January 2026].
- Engelke, P. (2020). *AI, Society, and Governance: An Introduction*. JSTOR. Available at: <https://www.jstor.org/stable/resrep29327> [Accessed 3 February 2026].
- Briggs and Dong (2025). *How Will AI Affect the Global Workforce?* [online] Goldmansachs.com. Available at: <https://www.goldmansachs.com/insights/articles/how-will-ai-affect-the-global-workforce> [Accessed 13 Oct. 2025].
- GOV.UK (2024). *UK Business Data Survey 2024*. Available at: <https://www.gov.uk/government/statistics/uk-business-data-survey-2024/uk-business-data-survey-2024> [Accessed 15 March 2026].
- Johnson-Roehr, S.N. (2023). Artificial Intelligence and Education: A Reading List. *JSTOR Daily*. Available at: <https://daily.jstor.org/artificial-intelligence-and-education-a-reading-list/> [Accessed 27 December 2025].

- Jung, C. and Desikan, B.S. (2024). *Scenarios for How Generative AI Could Affect Work in the Future*. JSTOR. Available at: <https://www.jstor.org/stable/resrep58388.7> [Accessed 2 October 2025].
- Rogers, E.M. (2003). *Diffusion of innovations*. 5th ed. New York: Simon and Schuster.
- Rossi, F. (2018). Building Trust in Artificial Intelligence. *Journal of International Affairs*, 72(1), pp. 127-134. Available at: <https://www.jstor.org/stable/26588348> [Accessed 20 December 2025].
- Sharps, S. (2024). The Impact of AI on the Labour Market. *Institute for Global Change*. Available at: <https://institute.global/insights/economic-prosperity/the-impact-of-ai-on-the-labour-market> [Accessed 12 March 2026].
- Shoukat, R. (2024). A Comparative Study of China, India and Pakistan. *Strategic Studies*, 44(1), pp. 89-108. doi:10.2307/48829133.
- Toh, A. and Ayoub, E. (2026). The Military's Use of AI, Explained. *Brennan Centre for Justice*. Available at: <https://www.brennancenter.org/our-work/research-reports/militarys-use-ai-explained> [Accessed 15 March 2026].
- Toloka Team (2023). Difference between AI, ML, LLM, and Generative AI. *Toloka AI Blog*. Available at: <https://toloka.ai/blog/difference-between-ai-ml-llm-and-generative-ai/> [Accessed 9 March 2026].
- United Nations (2024). Lethal Autonomous Weapons Systems. Report of the Secretary-General, A/79/88. Available at: un.org [Accessed 14 May 2026].

Appendix:

Survey Questions:

- Q1. Which of the following best describes your current role?
- Q2. How many years of professional experience do you have?
- Q3. How frequently do you use AI tools in your day-to-day work?
- Q4. Which AI tools or capabilities do you currently use at work?
- Q5. How has AI changed the way you personally work over the past two years?
- Q6. To what extent do you agree with the following statement? "AI will significantly reduce the number of jobs available in most sectors within the next 10 years."
- Q7. Which job types do you believe are most at risk of being replaced or significantly reduced by AI?
- Q8. Rank the following skills by importance for the future workforce. 1 = Most Important, 6 = Least Important: Communication and interpersonal skills / Critical thinking and problem-solving / Effective use of AI tools / Technical and coding skills / Adaptability and continuous learning / Creative and original thinking
- Q9. Do you believe companies should train employees to use AI as part of their role?
- Q10. Please briefly explain your reasoning for the answer above.
- Q11. In your opinion, which sectors will be most positively transformed by AI in the next decade?
- Q12. The use of AI in military and defence applications raises significant ethical questions. What is your view?
- Q13. How concerned are you about businesses relying on AI outputs without sufficient human oversight?
- Q14. Should schools and universities actively teach students how to use AI tools responsibly?
- Q15. To what extent do you agree with the following statement? "Using AI in education is comparable to the introduction of calculators — it should be embraced, not banned."

Q16. From your professional experience, what advice would you give to young people entering the workforce in an AI-driven world?

Key Results Summary:

Respondent Profile

- 57 total respondents
- 84% had more than 10 years of professional experience
- 84% already use AI at work at least sometimes
- 51% from the teaching sector
- Remaining respondents from business, finance, consulting, healthcare, law, engineering and sales

Q3 — AI Usage Frequency

- Daily: 10 respondents
- Often: 12 respondents
- Sometimes: 20 respondents
- Rarely: 12 respondents
- Never: 3 respondents
- 84% use AI at least sometimes at work

Q6 — Will AI significantly reduce jobs within 10 years?

- Average score: 3.3 out of 5
- 42% leaned towards agreement (scored 4 or 5)
- Only 4% strongly disagreed
- Not a single respondent thought no jobs would be affected at all

Q7 — Which job types are most at risk?

- Repetitive administrative roles: 55 respondents (96%)

- Financial analysis and reporting: 43 respondents
- Customer service and call centres: 39 respondents
- Software engineering and development: 35 respondents
- Legal and compliance: 28 respondents
- Healthcare diagnostics and reporting: 24 respondents
- Creative roles: 24 respondents
- Sales and account management: 19 respondents
- Senior leadership and strategy: 3 respondents
- 100% of respondents believed some jobs would be affected

Q8 — Skills ranking for the future workforce (lower average = ranked more important)

1. Communication and interpersonal skills — average 2.3
2. Critical thinking and problem-solving — average 2.5
3. Adaptability and continuous learning — average 2.7
4. Effective use of AI tools — average 3.1
5. Creative and original thinking — average 3.3
6. Technical and coding skills — average 3.5

Q9 — Should companies train employees to use AI?

- Yes definitely: 30 respondents
- Probably yes: 22 respondents
- Unsure: 4 respondents
- Probably not: 1 respondent
- No: 0 respondents
- 93% said yes — not a single respondent said no

Q11 — Sectors most positively transformed by AI

- Healthcare and Medicine: 38 respondents (67%)

- Transportation and Logistics: 27 respondents (47%)
- Financial Services: 24 respondents (42%)
- Legal and Compliance: 18 respondents
- Retail and E-Commerce: 14 respondents
- Education: 12 respondents
- Media and Creative Industries: 11 respondents
- Defence and Military: 8 respondents

Q13 — Concern about AI without human oversight

- Not concerned: 0 respondents
- Slightly concerned: 1 respondent
- Moderately concerned: 11 respondents
- Very concerned: 30 respondents
- Extremely concerned: 15 respondents
- 79% very or extremely concerned
- Nobody scored 1 or 2

Q14 — Should schools teach responsible AI use?

- Strongly yes: 25 respondents
- Yes: 29 respondents
- Unsure: 1 respondent
- No: 2 respondents
- 96% said schools should teach AI responsibly

Q15 — AI in education is like calculators — embrace not ban

- Strongly disagree: 0 respondents
- Disagree: 7 respondents
- Neutral: 20 respondents

- Agree: 18 respondents
 - Strongly agree: 12 respondents
-

Selected Open Text Responses:

On the future of work: *"AI will reward curiosity, not passivity. The most valuable people will be those who can ask good questions, spot weak answers, communicate clearly, adapt quickly and use human judgement."* — Finance Director

On military AI: *"Decisions involving life, death or targeting should not be delegated fully to AI. Accountability must remain with humans."* — Finance and Operations respondent

On military AI: *"The use of ethical AI is a huge problem that none of the LLMs currently focus on — but this has to become part of the conversation instead of revenue growth."* — Sales and Marketing respondent

On education: *"AI is a tool to be used, not a replacement for thought. Learn it, but don't rely on it."* — Teaching respondent

On advice for young people: *"Learn how to use AI well, but do not let it do your thinking for you."* — Multiple respondents